

Supportive — Week 1 Build Plan

Project context

Supportive is an Australian job directory for mental health and supportive services roles. The current codebase at this repo is a generic job board template ("JobBoard") with US tech jobs. Week 1's job is to transform it into a recognisably Australian MH job directory — nothing more, nothing less.

Read this file in full before starting any section. Do not skip sections or reorder them — later sections depend on earlier ones.

Working rules (apply to every session)

1. **Show diffs before committing.** Every change, show the user a diff and wait for approval.
2. **One section per session.** Each numbered section below is a separate Claude Code session. Do not try to do multiple sections in one turn.
3. **Commit at the end of each section** with a message like `Section N: <short description>`. Keep commits atomic.
4. **Do not scope-creep.** The "NOT in scope" list at the bottom of this file is binding — if something there feels tempting, stop and discuss with the user before proceeding.
5. **Update BUILD_NOTES.md as you go.** Create this file in the repo root if it doesn't exist. Log decisions, gotchas, and things to return to later.
6. **Ask before installing new dependencies.** If a section needs a package that isn't already installed, ask first.
7. **Prefer editing existing code over new abstractions.** This is a template being reshaped, not a greenfield rebuild.

The product we're building

- Name: **Supportive**
- Geography: **Australia only** — AUD, AU cities, AU states, DD/MM/YYYY dates
- Audience: candidates looking for mental health / AOD / peer work / NDIS / community services roles, and employers hiring for those roles
- Positioning: mission-aligned, not a generic job board
- Architecture: three-layer page model (editorial hubs → dense matrix pages → gated long tail), with a noindex threshold to prevent thin-content problems at scale

The 18-role taxonomy (LOCKED — do not suggest changes)

This is the spine of the site. Every route, filter, editorial page, and feed integration references these roles. They are locked for Week 1. If a change is proposed, route the proposal to the user — do not change them unilaterally.

Clinical

1. Psychologist ([psychologist](#))
2. Clinical Psychologist ([clinical-psychologist](#))
3. Psychiatrist ([psychiatrist](#))
4. Mental Health Nurse ([mental-health-nurse](#))
5. Mental Health Occupational Therapist ([mental-health-occupational-therapist](#))

Allied health

6. Mental Health Social Worker ([mental-health-social-worker](#))
7. Counsellor ([counsellor](#))
8. Psychotherapist ([psychotherapist](#))
9. Art / Music / Play Therapist ([creative-therapist](#))
10. Dietitian (Eating Disorders) ([eating-disorder-dietitian](#))

Community / NDIS

11. Peer Support Worker ([peer-support-worker](#))
12. Lived Experience Worker ([lived-experience-worker](#))
13. Support Worker (Mental Health) ([mental-health-support-worker](#))
14. AOD Worker ([aod-worker](#))
15. Youth Mental Health Worker ([youth-mental-health-worker](#))
16. Case Manager / Care Coordinator ([case-manager](#))

Leadership & specialist

17. Clinical Lead / Team Leader ([clinical-lead](#))
18. Mental Health Educator / Trainer ([mental-health-educator](#))

Category slugs: [clinical](#), [allied-health](#), [community-ndis](#), [leadership](#).

AU locations seed list

Cities: Sydney, Melbourne, Brisbane, Perth, Adelaide, Canberra, Hobart, Darwin, Newcastle, Wollongong, Gold Coast, Geelong. Plus "Remote (Australia)" and "Hybrid".

States (enum): NSW, VIC, QLD, WA, SA, ACT, TAS, NT.

Section 0 — Pre-flight audit

Goal: map the codebase before changing anything.

Tasks:

1. Summarise the tech stack: framework, styling, DB/data layer, auth, deployment target.
2. Map the folder structure at the top level and for any folder deeper than one level that looks load-bearing.
3. Identify where the job data currently lives: hardcoded? seeded DB? mock JSON? Supabase? Prisma schema? Report the exact file(s).
4. Find every file that mentions `JobBoard`, `Seattle`, `San Francisco`, `USD`, category names like `Engineering`, `DevOps`, `Product`, `Design`, `Data Science`. Output as a list grouped by concern.
5. Identify where theme/colors/brand tokens are defined (Tailwind config? CSS vars? shadcn theme file? a `theme.ts`?).
6. Create `BUILD_NOTES.md` in the repo root with the headings `Decisions`, `Gotchas`, `To revisit`, `Architecture notes`. Put the pre-flight summary under `Architecture notes`.

Deliverable: a written summary back to the user covering the six tasks above. No code changes.

Commit: none yet. This is read-only.

Section 1 — Brand and identity swap

Goal: rename the site from JobBoard to Supportive and reshape the hero copy for mental health.

Tasks:

1. Find-and-replace `JobBoard` → `Supportive` across the entire codebase. Check component files, page titles, meta tags, footer, alt text, aria labels.
2. Update the hero H1: `Find Your Dream Job` → `Find your next role in mental health`.

3. Update the hero subcopy to: Roles across clinical mental health, AOD, peer work, NDIS, and community services. Posted by mission-aligned employers.
4. Update the tagline/meta-description wherever it lives to: Mental health and supportive services careers, Australia-wide.
5. Update footer copyright year to the current year and brand to Supportive.
6. Replace the logo with a placeholder — just an S monogram in a square with the primary theme color is fine. Do not rabbit-hole on branding.
7. Update the favicon to match (a simple generated SVG favicon with S is fine).
8. Update the "For Job Seekers" / "For Employers" / "Company" footer groupings — text is fine as-is, but if any links reference tech-specific copy, update them to be role-neutral.

Acceptance criteria:

- Loading the site shows no JobBoard string anywhere.
- Hero H1 reads Find your next role in mental health.
- Meta title tag and meta description both mention Supportive and mental health / Australia.

Deliverable: diff, then commit with message Section 1: Brand and identity swap to Supportive.

Section 2 — Australia localisation

Goal: make every surface read as Australian.

Tasks:

1. Change all currency references from USD to AUD. The \$ symbol stays the same — do not use A\$ unless a specific component explicitly needs to disambiguate.
2. Remove every US city from seed data, sample jobs, and dummy components. If a placeholder location is needed, use Sydney, NSW.
3. Create an AU locations seed module (location name these in a sensible place — likely lib/constants/locations.ts or similar depending on the stack). Include the 12 cities plus Remote (Australia) and Hybrid listed above. Include a state field for each.
4. Create an AU states enum/constant in the same style: NSW, VIC, QLD, WA, SA, ACT, TAS, NT.
5. Convert date formatting anywhere it appears from MM/DD/YYYY to DD/MM/YYYY. If relative formatting is already in use (2 days ago), leave it.

6. On the homepage stats block (500+ Active Jobs / 100+ Companies / 10+ Categories / 2,000+ Hired): the current numbers are fabricated. Either (a) hide the entire stats block behind a feature flag defaulting to off until we have real numbers, or (b) change the framing to honest pre-launch copy — e.g. Launching with: 18 role categories, 8 states, mission-aligned employers. Ask the user which they'd prefer before committing.
7. Check for any i18n/locale strings and set the default locale to en-AU.

Acceptance criteria:

- No US cities appear anywhere visible to a visitor.
- All salary ranges show as AUD.
- Stats block is either hidden or shows honest pre-launch copy.

Deliverable: diff, then commit with message Section 2: Australia localisation.

Section 3 — Taxonomy swap

Goal: replace the tech category taxonomy with the 18-role mental health taxonomy.

This is the most important section. Everything downstream depends on it.

Tasks:

1. Replace the current category enum / data structure with the 18 roles listed in the "18-role taxonomy" section at the top of this file. Preserve both the display name and the slug for each.
2. Add a parent category grouping using the four category slugs: clinical, allied-health, community-ndis, leadership.
3. Update the filter UI on the jobs list page to show the 18 roles, grouped visually under the four category headings.
4. Update the homepage "Categories" stat (if kept) to reflect 18 roles in 4 groups.
5. Replace all six existing sample jobs on the homepage with these six Australian MH roles. Keep the same card structure, just swap the data:

Role	Employer	Location	Type	Salary (AUD)
Peer Support Worker	headspace	Melbourne, VIC	Part-time	\$70,000–\$85,000
Clinical Psychologist	Black Dog Institute	Sydney, NSW	Full-time	\$115,000–\$140,000
Mental Health Social Worker	Neami National	Brisbane, QLD	Full-time	\$90,000–\$105,000
AOD Worker	Odyssey House	Melbourne, VIC	Full-time	\$80,000–\$95,000
Mental Health Nurse	NSW Health	Newcastle, NSW	Full-time	\$100,000–\$120,000
Youth Mental Health Worker	ReachOut	Remote (Australia)	Part-time	\$75,000–\$90,000

6. Mark the first three in the list as `featured: true` so the homepage featured section still has content.
7. On individual job detail pages (`/jobs/[id]`), verify that all six sample jobs render without errors and show the correct role/category badge.
8. If there's a DB schema (Prisma, Drizzle, Supabase migrations, etc.): update the schema to support the new taxonomy. Do NOT drop existing data without asking — if there's real data, discuss first. For a template with seed data only, reseeding is fine.

Acceptance criteria:

- Filters show 18 roles grouped by 4 categories.
- Homepage displays 6 Australian MH sample jobs with AUD salaries.
- Zero references to `Engineering`, `DevOps`, `Product`, `Design`, `Data Science` anywhere in the codebase.

Deliverable: diff, schema migration (if applicable), then commit with message `Section 3: 18-role MH taxonomy`.

Section 4 — URL structure and routing

Goal: establish the three-layer page architecture routing skeleton. Do not build out full page content — just make the routes work with placeholders.

Routes to create:

None	
/	Homepage (exists)
/jobs	All jobs, filterable (exists)
/jobs/[id]	Individual job (exists)
/roles	Index of all 18 roles (new)
/roles/[role-slug]	Editorial hub per role (new, Layer 1)
/roles/[role-slug]/[location-slug]	Role × location matrix (new, Layer 2)
/locations	Index of all locations (new)
/locations/[location-slug]	Location hub (new)
/employers	Index of employers (new)
/employers/[employer-slug]	Employer hub (new, Layer 1)
/specialties	Index of specialty areas (new)
/specialties/[specialty-slug]	Specialty hub (new, Layer 1)
/blog	Blog index (exists)
/blog/[slug]	Blog post (exists)
/about	About page (exists)

Tasks:

1. Create each new route file with a minimal placeholder page component.
2. Each placeholder should:
 - Pull the relevant name from the URL (role name, location name, employer name, etc.)
 - Display that name as the H1
 - Show a short **Coming soon – this hub is being built** message
 - Set a proper `<title>` and meta description using the pulled name
 - Emit a canonical link tag matching the current URL
3. Generate the specialty hub list to start with: AOD / Addiction, Youth Mental Health, Perinatal Mental Health, Eating Disorders, Suicide Prevention, Trauma, Forensic Mental Health, Aboriginal & Torres Strait Islander Mental Health, LGBTQIA+ Mental Health, Rural & Remote.
4. Generate the index pages (`/roles`, `/locations`, `/employers`, `/specialties`) as simple alphabetical lists linking to each detail page.

5. Update the sitemap generator to include all new routes. The sitemap must include every role slug, every location slug, every role×location combination, every employer slug, and every specialty slug.
6. Add navigation updates: the main nav should link to `/jobs` and `/roles` at minimum. Footer should link to `/roles`, `/locations`, `/employers`, `/specialties`.

Acceptance criteria:

- Navigating to `/roles/psychologist` shows "Psychologist" as H1 with Coming Soon message.
- Navigating to `/roles/psychologist/sydney` shows "Psychologist jobs in Sydney" as H1 with Coming Soon message.
- Sitemap XML includes all generated routes.
- No 404s on any of the routes listed above.

Deliverable: diff, then commit with message `Section 4: Three-layer routing skeleton`.

Section 5 — Noindex threshold logic

Goal: build server-side logic that prevents thin pages from being indexed by Google until they meet a quality threshold.

The rule — a page is indexable only if ALL three are true:

1. At least 3 active job listings on the page, OR at least 5 historical listings displayed as "recent roles"
2. At least 300 words of unique content on the page (excluding navigation, footer, and listing card text)
3. A meta description that is not auto-generated boilerplate (i.e. it includes unique content beyond just `Jobs in X / Y roles in Z`)

Where to apply the rule:

- `/roles/[role-slug]/[location-slug]` — yes, apply
- Any future Layer 3 routes — yes, apply (leave hook in place)
- `/roles/[role-slug]` — NO, always indexable (Layer 1 editorial hub)
- `/employers/[employer-slug]` — NO, always indexable
- `/specialties/[specialty-slug]` — NO, always indexable
- `/jobs/[id]` — NO, always indexable (individual job pages)

- `/jobs`, `/`, static pages — NO, always indexable

Tasks:

1. Build a server-side helper function `shouldNoindex(pageData): boolean` that takes a page's data (listings count, unique content word count, meta description) and returns true/false per the three conditions.
2. Wire the helper into the relevant dynamic routes so that if `shouldNoindex()` returns true, the page emits `<meta name="robots" content="noindex, follow" />` in the head. If false, emit no robots meta tag (default = indexable).
3. Also update the sitemap generator: gated (`noindex`) pages must NOT appear in the sitemap. Re-running the sitemap generator should only include currently-indexable pages.
4. Build a dev-only `/debug/index-status` page (protected by env check — only accessible when `NODE_ENV !== 'production'` or behind an admin flag). It should list every dynamic page and show:
 - URL
 - Current listings count
 - Current unique content word count
 - Whether the meta description is unique
 - Overall indexability status (green/amber/red)
5. Add unit tests for the `shouldNoindex` helper covering: zero listings + short content (`noindex`), 5 listings + short content (`noindex`, fails word count), 5 listings + 400 words + generic meta (`noindex`, fails meta check), 3 listings + 400 words + unique meta (`indexable`).

Acceptance criteria:

- Visiting `/roles/psychologist/darwin` with zero listings shows `noindex` meta tag in page source.
- Visiting `/roles/psychologist` shows no `noindex` meta tag (Layer 1 always indexable).
- `/debug/index-status` is accessible in dev and shows the full table.
- Sitemap excludes `noindex` pages.
- Unit tests pass.

Deliverable: diff, helper function, integration, debug page, tests, then commit with message `Section 5: Noindex threshold system.`

Section 6 — JSON-LD JobPosting structured data

Goal: emit Google JobPosting schema on individual job pages so the site becomes eligible for Google for Jobs.

Reference: <https://developers.google.com/search/docs/appearance/structured-data/job-posting>

Tasks:

1. On `/jobs/[id]`, emit a `<script type="application/ld+json">` block in the `<head>` with the `JobPosting` schema.
2. Required fields (Google spec):
 - `@context`: `https://schema.org/`
 - `@type`: `JobPosting`
 - `title`: job title
 - `description`: job description (HTML-encoded)
 - `datePosted`: ISO 8601
 - `validThrough`: ISO 8601 — default to 30 days after `datePosted` if not set
 - `employmentType`: map from the site's enum to Google's: `FULL_TIME`, `PART_TIME`, `CONTRACTOR`, `TEMPORARY`, `INTERN`, `VOLUNTEER`, `PER_DIEM`, `OTHER`
 - `hiringOrganization`: { `"@type"`: `"Organization"`, `"name"`: `employerName`, `"sameAs"`: `employerUrl` (optional) }
 - `jobLocation`: { `"@type"`: `"Place"`, `"address"`: { `"@type"`: `"PostalAddress"`, `"addressLocality"`: `city`, `"addressRegion"`: `state`, `"addressCountry"`: `"AU"` } }
 - `baseSalary`: { `"@type"`: `"MonetaryAmount"`, `"currency"`: `"AUD"`, `"value"`: { `"@type"`: `"QuantitativeValue"`, `"minValue"`: `min`, `"maxValue"`: `max`, `"unitText"`: `"YEAR"` } }
3. For remote roles, also include `jobLocationType`: `"TELECOMMUTE"` and `applicantLocationRequirements`: { `"@type"`: `"Country"`, `"name"`: `"Australia"` }.
4. Build a helper `buildJobPostingSchema(job): object` that constructs the JSON-LD object from a job record. Reuse across any page that needs it.
5. Add a note in `BUILD_NOTES.md` under `To revisit`: "Eventually also emit `ItemList` JSON-LD on `/jobs` index and role/location hubs for better discovery."

Acceptance criteria:

- Viewing a job page's source shows a valid JSON-LD block.

- Google's Rich Results Test tool (<https://search.google.com/test/rich-results>) validates the output without errors.
- Remote jobs include the telecommute fields.

Deliverable: diff, helper function, test with Google's Rich Results Test, then commit with message `Section 6: JobPosting JSON-LD`.

Section 7 — Deploy and walkthrough

Goal: deploy everything, then do a full walkthrough to catch anything that feels off before Week 2.

Tasks:

1. Run `git log` and confirm all six previous sections are committed cleanly.
2. Verify the build passes locally (`npm run build` or the equivalent).
3. Confirm the Vercel deployment succeeds.
4. Walk through the live site as two personas and report findings:
 - **Candidate persona:** "I'm a peer worker in Melbourne looking for part-time work in youth mental health." Can they find a job? Can they filter sensibly? Does the UX make sense?
 - **Employer persona:** "I run a 30-person AOD service in Brisbane and I'm thinking of posting three roles." Does the site communicate what it is? Is it obvious how to post? Is the register/login flow coherent?
5. List every friction point, broken link, "this looks wrong" moment, or missing piece. Do NOT fix them — just document in `BUILD_NOTES.md` under `To revisit` for Week 2 prioritisation.
6. Summarise: does the site pass the Week 1 definition of done? (See below.)

Definition of done — all three must be "yes":

1. Is this an Australian job board? (AUD, AU cities, AU spelling, AU locale throughout.)
2. Is this focused on mental health and supportive services? (18 roles, MH-specific hero copy, MH sample jobs, MH employers.)
3. Would a candidate or employer in this space bother to come back? (Site feels like it was built for them, not adapted from a tech template.)

Deliverable: deployment, walkthrough notes in `BUILD_NOTES.md`, final summary back to the user.

Explicitly NOT in Week 1 scope

If you find yourself reaching for any of these, stop and route to the user:

- Real job feed integration (Adzuna, Greenhouse, Lever, Workday, PageUp) — Week 2
- Full editorial hub content writing — Week 3 (just the first 5 roles)
- Employer dashboard / job posting UI — Week 4+
- Payments, Stripe, subscription tiers, featured-listing upsells — post-launch
- Matching wizard, AI recommendations, semantic search — post-launch
- Email notifications, job alert signups — post-launch
- Auth flow polish beyond the existing template — only if actively broken
- Analytics beyond basic Vercel Analytics — post-launch
- Third-party integrations (Slack, Typeform, etc.) — post-launch
- Any ML or AI features — post-launch

If something feels like it should happen now but is on this list, it probably shouldn't. Surface it to the user rather than acting on it.

Cross-cutting reminders

- The 18-role taxonomy is locked. Do not propose changes.
- The AU locations list is locked for Week 1. Do not expand.
- The noindex threshold rule is locked. If exceptions seem needed, discuss first.
- Sentence case for all UI copy (including buttons, labels, headings). No Title Case.
- Australian English spelling throughout (**organisation**, **specialise**, **colour**).
- Never say "lifetime access", "guaranteed", or any marketing superlative. This is a directory.
- When in doubt, ask the user. Cheap to clarify, expensive to redo.